Exact Orthogonal-CI Contraction for Arbitrary Selected Valence-Bond Structure Spaces

Theory and implementation note

16 October 2026

Abstract

An arbitrary selected valence-bond (VB) structure space is generally not invariant under orbital orthogonalization. This does not prevent exact evaluation of its Hamiltonian and overlap actions through an orthogonal configuration-interaction (CI) representation. The transformed wave function must be retained, but its complete spin-string product array need not be formed. We derive an exact source-factor representation, apply the Slater–Condon one-body generators separately to the two spin factors, and contract only the original products needed for the final structure-space projection. The resulting equations preserve the selected variational space without a closure assumption or intermediate amplitude truncation. Source and channel tiling bound the numerical workspace. Optional spin-exchange reduction and causal pruning of the exterior adjoint are derived separately. The scope is the Hamiltonian/overlap action, its active-integral adjoint, directional integral adjoints, oriented residual-generator adjoints, and selected-structure coefficient-response contractions.

1 The selected VB space

Let n be the number of active spatial orbitals, with n_α and n_β active electrons of each spin. The complete one-spin string spaces have dimensions

$$U_\alpha = \binom{n}{n_\alpha}, \qquad U_\beta = \binom{n}{n_\beta}. \tag{1}$$

An index a identifies an α string and b a β string. For a selected VB structure I, let $w_{ab,I}$ be its signed expansion coefficient on the corresponding spin-string product, with all normalization conventions included. For structure coefficients $\mathbf{C}$, define

$$X = \mathcal{W}\mathbf{C}, \qquad X_{ab} = \sum_I w_{ab,I} C_I. \tag{2}$$

Here $\mathcal{W}$ is a sparse linear map from structure coefficients to a matrix of spin-string product coefficients. The matrix notation does not prescribe a dense allocation: the map can be stored as sparse incidences between structures and touched string products. It introduces no determinant-pair traversal or separately optimized determinant coefficient vector.

The original support and its two marginal index sets are

$$\Omega = \{(a,b) : w_{ab,I} \neq 0 \text{ for at least one selected } I\}, \qquad t = |\Omega|, \tag{3}$$

$$A_\Omega = \{a : \exists b,\ (a,b) \in \Omega\}, \qquad J = \{b : \exists a,\ (a,b) \in \Omega\}, \qquad r = |J|. \tag{4}$$

Support is determined by the selected structures, not by numerical smallness or cancellations in a particular Davidson vector. Using the Frobenius inner product, the adjoint structure map is

$$(\mathcal{W}^* Z)_I = \sum_{(a,b)\in\Omega} w_{ab,I} Z_{ab}. \tag{5}$$

Thus only entries in Ω are needed at the *end* of a complete action. Nothing in this statement permits dropping intermediate entries outside Ω.

2 Orthogonalization without a closure assumption

At a fixed orbital point, write the nonorthogonal active orbitals as $\phi = \chi R$, where χ is an orthonormal basis of the same active one-particle space. For example, if O is their positive-definite overlap,

$$O = R^{\mathsf{T}} R, \qquad \chi = \phi R^{-1}. \tag{6}$$

This is an evaluation-coordinate change, not a change to the physical localization constraints or to the orbital optimization variables. Multilinearity of antisymmetrized products gives the exterior maps

$$T_\sigma = \wedge^{n_\sigma} R, \qquad (T_\sigma)_{ai} = \det R[a, i], \qquad \sigma \in \{\alpha, \beta\}, \tag{7}$$

where the row and column subsets contain n_σ orbitals with a fixed ordering convention. The orthogonal-basis coefficient matrix is

$$Y = \mathcal{T} X = T_\alpha X T_\beta^{\mathsf{T}}, \qquad \mathcal{T}^* Z = T_\alpha^{\mathsf{T}} Z T_\beta. \tag{8}$$

The minor matrices T_σ are mathematical definitions; the implementation applies elementary exterior transforms without assembling them densely.

Let $\mathcal{H}$ be the active Hamiltonian action in the orthogonal spin-string product basis. The selected structure Hamiltonian and overlap are exactly

$$\boxed{H = \mathcal{W}^* \mathcal{T}^* \mathcal{H} \mathcal{T} \mathcal{W}, \qquad M = \mathcal{W}^* \mathcal{T}^* \mathcal{T} \mathcal{W}.} \tag{9}$$

An orbital-dependent scalar inactive-plus-nuclear energy E_0 contributes $E_0 M$ to the total structure Hamiltonian and is separate from the active action H in this note. The variational energy is

$$E = E_0 + \frac{\mathbf{C}^{\mathsf{T}} H \mathbf{C}}{\mathbf{C}^{\mathsf{T}} M \mathbf{C}}, \tag{10}$$

with H denoting the active contribution of Eq. (9).

Proposition 1 (Exact selected-space restriction)**.** *For any selected structure map $\mathcal{W}$, Eq.* (9) *gives the exact selected-space Hamiltonian and overlap, whether or not $\mathcal{T}$* $\operatorname{ran} \mathcal{W} \subseteq \operatorname{ran} \mathcal{W}$.

Proof. For arbitrary structure vectors $\mathbf{u}$ and $\mathbf{v}$, the two physical states have orthogonal-basis amplitudes $\mathcal{T}\mathcal{W}\mathbf{u}$ and $\mathcal{T}\mathcal{W}\mathbf{v}$. Their Hamiltonian matrix element is

$$\langle \mathcal{T}\mathcal{W}\mathbf{u}, \mathcal{H}(\mathcal{T}\mathcal{W}\mathbf{v}) \rangle_{\mathrm{F}} = \mathbf{u}^{\mathsf{T}} \big(\mathcal{W}^* \mathcal{T}^* \mathcal{H} \mathcal{T} \mathcal{W} \big) \mathbf{v}. \tag{11}$$

Replacing $\mathcal{H}$ by the identity gives the overlap. No intermediate restriction to $\operatorname{ran} \mathcal{W}$ was inserted. □

The nonclosure issue arises only if a different operation is used, for example by zeroing transformed amplitudes outside the original selected support before applying $\mathcal{H}$. That operation changes the wave function. Equation (9) does not do so. Davidson vectors and Ritz coefficients remain in the selected VB structure space; the orthogonal representation is used internally to evaluate their images. For linearly independent selected structures, their eigenproblem remains

$$H\mathbf{C} = E_{\mathrm{act}} M \mathbf{C}. \tag{12}$$

The adjoint $\mathcal{W}^*$ contracts against the selected test structures; it is not an inverse coefficient transform. The nonidentity overlap M is retained, so this procedure does not replace the selected VB problem by an unconstrained orthogonal-CI eigenproblem.

2.1 A minimal nonclosed example

Consider a two-by-two coefficient matrix with the single selected product $X = e_2 e_2^{\mathsf{T}}$, and let

$$T_\alpha = \begin{pmatrix} 1 & a \\ 0 & 1 \end{pmatrix}, \qquad T_\beta = \begin{pmatrix} 1 & b \\ 0 & 1 \end{pmatrix}. \tag{13}$$

This is a linear-algebra example of the coefficient map, not an assertion that a single spin product is a spin-adapted VB structure. Its transformed amplitudes are

$$Y = \begin{pmatrix} ab & a \\ b & 1 \end{pmatrix} = \begin{pmatrix} a \\ 1 \end{pmatrix} \begin{pmatrix} b \\ 1 \end{pmatrix}^{\mathsf{T}}. \tag{14}$$

For nonzero a and b, all four products are present, so the original one-dimensional selected space is not closed. Nevertheless, two vectors represent the complete Y exactly. The overlap is $(1 + a^2)(1 + b^2)$, not the value 1 obtained by retaining only Y_{22}. An exact Hamiltonian contraction must likewise use all transformed amplitudes, explicitly or implicitly. The factorized route below retains them implicitly.

3 Exact source factors

3.1 Column decomposition

Order the original touched β columns as $J = (j_1, \ldots, j_r)$, and set

$$A_{:\ell} = X_{:j_\ell}, \qquad B_{:\ell} = e_{j_\ell}, \qquad X = AB^{\mathsf{T}}. \tag{15}$$

This is an identity for every input vector. It is not a truncated singular value decomposition, and r is not a fitted rank. Transforming only the two factors gives

$$L = T_\alpha A, \qquad F = T_\beta B, \qquad \mathcal{T}X = LF^{\mathsf{T}}. \tag{16}$$

The factors have dimensions $U_\alpha \times r$ and $U_\beta \times r$. The complete transformed $U_\alpha \times U_\beta$ matrix need not be formed, even if every one of its entries is nonzero.

The distinction between r and the number of structures matters. A single covalent structure with m disjoint singlet pairs generally touches 2^m original source columns. Several structures can share most of those columns, and a broad selected space can touch all U_β columns. Thus rank $X \le r$ is an upper bound, not an identification of r with one, with rank X, or with the number of structures.

3.2 A minimum mixed row/column cover

Grouping every source on the β side is not always economical. View the original input support as a bipartite graph: vertices are the touched α and β strings, and each supported product (a, b) is one edge. Choose a vertex cover $C_\alpha \cup C_\beta$, so that every edge touches at least one chosen vertex. Assign an edge to its α endpoint if that endpoint is in C_α; otherwise assign it to its covered β endpoint. With this disjoint assignment,

$$X = \sum_{a \in C_\alpha} e_a u_a^{\mathsf{T}} + \sum_{b \in C_\beta} v_b e_b^{\mathsf{T}} = AB^{\mathsf{T}}. \tag{17}$$

Here $(u_a)_b = X_{ab}$ only for edges assigned to row a, and $(v_b)_a = X_{ab}$ only for edges assigned to column b; the remaining entries are zero. Thus no coefficient is counted twice, including edges whose two endpoints are both covered. Each chosen vertex defines one coordinate-anchored outer product.

Let $\nu(\Omega)$ be the maximum matching size, i.e., the largest number of edges with no shared endpoints. A minimum cover has cardinality

$$r_* = |C_\alpha| + |C_\beta| = \nu(\Omega) \le \min(|A_\Omega|, |J|). \tag{18}$$

For completeness, the cover can be certified directly from a maximum matching. Start at unmatched α vertices and follow alternating unmatched/matched edges. If the reached vertex sets are Z_α and Z_β, then $(A_\Omega \setminus Z_\alpha) \cup Z_\beta$ covers every edge and contains exactly one endpoint per matched edge. Any cover must contain at least one endpoint of each disjoint matched edge, establishing minimality. The implementation uses layered augmenting paths, with $O(t\sqrt{U_\alpha + U_\beta})$ worst-case planning work and linear graph storage. The plan is independent of the numerical Davidson vector.

This is a minimal factor count for *support-only coordinate-anchored* decompositions, not a numerical rank claim. In particular, rank $X \le \nu(\Omega)$; additional coefficient relations can lower the actual rank further. A diagonal support needs one source per edge. In contrast, the support

$$\{(a_0, b_j) : 1 \le j \le m\} \cup \{(a_i, b_0) : 1 \le i \le m\} \tag{19}$$

needs only two mixed sources, although either a pure row or pure column decomposition needs $m + 1$ sources. This reduces the factor dimension, rather than enlarging a tile or imposing an amplitude threshold.

All subsequent Hamiltonian and sampling equations apply unchanged, with $r = r_*$ and the factors of Eq. (17). When a common exchange parity is used, construct the graph from the *representative input* support of Section 6.1; the final projection still uses the complete original Ω. If the original column cover is already minimal, its ordering may be retained to avoid changing established summation order.

4 Hamiltonian action on the factors

4.1 One-body channels and the complete active Hamiltonian

Use real spatial orbitals and real, spin-independent molecular integrals. Let $e_{pq} = \sum_\sigma a^\dagger_{p\sigma} a_{q\sigma}$, and use the chemists' notation $(pq|rs)$ for two-electron integrals. The active Hamiltonian is

$$\hat{H} = \sum_{pq} h_{pq} e_{pq} + \frac{1}{2} \sum_{pqrs} (pq|rs)\big(e_{pq} e_{rs} - \delta_{qr} e_{ps}\big). \tag{20}$$

Here h includes the effective one-electron contribution of the inactive orbitals. Define

$$f_{ps} = h_{ps} - \frac{1}{2} \sum_q (pq|qs), \qquad N = n_\alpha + n_\beta > 0. \tag{21}$$

Packed pair labels $P = (p, q)$ have $p \ge q$, so there are $d = n(n+1)/2$ channels. Define the symmetric generators

$$\hat{E}_{\sigma P} = \begin{cases} a^\dagger_{p\sigma} a_{p\sigma}, & p = q, \\ a^\dagger_{p\sigma} a_{q\sigma} + a^\dagger_{q\sigma} a_{p\sigma}, & p > q, \end{cases} \qquad \hat{E}_P = \hat{E}_{\alpha P} + \hat{E}_{\beta P}. \tag{22}$$

Let $f_P = f_{pq}$, $V_{PQ} = (pq|rs)$ for $Q = (r, s)$, and $d_P = 1$ for diagonal pairs and zero otherwise. A symmetric kernel containing the *entire* active Hamiltonian is

$$K_{PQ} = \frac{1}{2} V_{PQ} + \frac{f_P d_Q + d_P f_Q}{2N}, \qquad \hat{H} = \sum_{PQ} K_{PQ} \hat{E}_P \hat{E}_Q. \tag{23}$$

Indeed, $\sum_P d_P \hat{E}_P = \hat{N} = N\hat{I}$ on the fixed-electron space. The second term in K therefore gives $\sum_P f_P \hat{E}_P$; the first term supplies the quadratic part of Eq. (20). The usual integral permutation symmetries account for both ordered off-diagonal orbital pairs through $\hat{E}_P$; no extra off-diagonal multiplicity is needed in K.

Let $E_{\sigma P}$ be the matrix of $\hat{E}_{\sigma P}$ on its one-spin string space. The Slater–Condon rules make this operation a sparse signed occupation/replacement scatter, rather than a dense string pair contraction. Its action on a product matrix is

$$\mathcal{E}_P(Z) = E_{\alpha P} Z + Z E_{\beta P}^{\mathsf{T}}, \qquad \mathcal{H} = \sum_{PQ} K_{PQ} \mathcal{E}_P \mathcal{E}_Q. \tag{24}$$

These identities require a complete carrier for each fixed-spin sector, or an exactly equivalent representation of all required intermediates. Arbitrary *selected VB structure spaces* need not be complete. The two statements are different.

4.2 Generated channel sweep and its adjoint

Let S contain the exact support of the transformed selected input and $I = \bigcup_P \mathrm{supp}(\mathcal{E}_P S)$ its one-generator closure. For a response L and state Y, define the channel panels

$$D^L_{zP} = (\mathcal{E}_P L)_z, \qquad D^Y_{zP} = (\mathcal{E}_P Y)_z, \qquad z \in I. \tag{25}$$

The same generated paths supply the Hamiltonian image and its integral cotangent:

$$\sigma_L = \mathcal{E}^*(D^L K^{\mathsf{T}}), \qquad G_K = (D^L)^{\mathsf{T}} D^Y. \tag{26}$$

Here $\mathcal{E}^*$ is the signed occupation/replacement gather adjoint to the channel scatter. The Hamiltonian image is not truncated back to S before any further consumer; its exterior-transform and structure adjoints perform the exact final projection. The supplied support must contain both inputs. For $L = Y$ the state scatter is the response scatter and is reused directly. For distinct inputs, read the state into the consumed input halo and place D^Y in the slab that subsequently holds $D^L K^{\mathsf{T}}$. After accumulating G_K, that slab is reused for sigma. Thus a second input requires no extra product array, channel slab, or worker-local tensor.

For two inputs with the same certified spin-exchange parity, the channel sum may use one representative per exchange orbit. Its multiplicity is two off the exchange diagonal and one on the diagonal; apply this weight to only one Gram operand. This gives the full integral cotangent, while the emitted sigma remains a projected-action representative until its known parity projection is applied. Uncertified or unequal spin spaces use the complete generated channel support, without an exchange quotient.

With $m = |I|$ generated channel targets and b right-hand sides, the dense channel-kernel and Gram contractions cost $O(mbd^2)$ each, in addition to the sparse scatters, readers and final projections. The complete-carrier order instead uses $m = U_\alpha U_\beta$. This optimization avoids absent targets and repeated scatter; it does not lower the power of the orbital count in the exact channel kernel. The returned integral cotangent needs $O(d^2)$ storage, separately from the fixed shared numerical slab budget. No higher-order density tensor is formed, and no closure of the selected VB structure space is assumed.

4.3 Joint orbital-direction action

At a fixed orthogonal evaluation basis the signed spin-string connections do not depend on the integrals. Their orbital dependence has already been transferred to Y and the channel kernel K. For an orbital direction let $Z = \delta Y$ and $\dot{K} = \delta K$, and define $A = \mathcal{E} Z$ and $B = \mathcal{E} Y$. Differentiating the Hamiltonian image gives

$$\delta(\mathcal{H} Y) = \mathcal{H} Z + \delta\mathcal{H} Y = \mathcal{E}^* \left(A K^{\mathsf{T}} + B \dot{K}^{\mathsf{T}} \right). \tag{27}$$

Here $\dot{K}$ includes the transformed one-electron and two-electron directions in Eq. (23); it is not merely the two-electron direction. The integral cotangent required by the orbital Hessian pullback is obtained from the same two scatters:

$$J_K(Z, Y) = A^{\mathsf{T}} B. \tag{28}$$

State weights and the factor of two for real-state energy derivatives are applied after this bilinear contraction. They must not weight B in Eq. (27) itself.

The equality in Eq. (27) eliminates the second connection/halo traversal, the repeated scatter of Y, and a separate physical image gather. It also replaces the two orbital-generator cotangents of $\mathcal{H}Z$ and $\delta\mathcal{H}Y$ by the cotangent of their sum, by linearity in the left argument. Both kernel multiplications remain necessary; neither the $O(mbd^2)$ dense channel-kernel scaling nor the selected variational space changes. In particular, this is not a low-rank approximation or a relaxation of the Hessian accuracy.

Three shared channel slabs suffice: retain A and B, form $A^\mathsf{T}B$, write $C = AK^\mathsf{T}$ into the third slab, overwrite the now-consumed A by $B\dot{K}^\mathsf{T}$, and add it to C. The same numerical budget is retained by shortening the generated target batch when required. The extra slab therefore does not imply an increased cache budget or an additional spin-product array. Smaller batches can increase metadata overhead; a single symbolic traversal does not guarantee a factor-of-two speedup.

For a certified spin-exchange parity, let D be the diagonal matrix of exchange-orbit multiplicities on generated target rows. The integral cotangent is $A^\mathsf{T}DB$. The multiplicity must be removed from the state operand before forming $B\dot{K}^\mathsf{T}$; the projected-image multiplicity is then applied exactly once to the combined output. Unequal spin spaces or uncertified parity use $D = I$.

The raw finite-orbital direction still contains the independent left exterior-transform term:

$$\delta(H\mathbf{C}) = \mathcal{W}^* \left[(\delta\mathcal{T})^*\mathcal{H}Y + \mathcal{T}^*\{\mathcal{H}Z + \delta\mathcal{H}Y\}\right] \quad (\delta\mathbf{C} = 0). \tag{29}$$

This term is not eliminated by channel fusion. When accepted images are not resident, a second right-hand side carries Y while the state reader has a structural zero in that slot. The joint outputs are then $(\mathcal{H}Y, \mathcal{H}Z + \delta\mathcal{H}Y)$, supplying the left term without an accepted-point replay outside the shared sweep. Optional diagnostics may gather $\delta\mathcal{H}Y$ separately from the already-computed overwritten slab; the production path needs only the combined image.

4.4 Triangular directions preserve the generated support

The support S used above is the downward closure of the original touched spin products under the upper-triangular exterior transform, not the selected VB structure subspace itself. Write P_S for its coordinate projector. Every e_{pq} with $p \le q$ either keeps an occupation unchanged or replaces an occupied index by a lower index in one spin sector. Consequently,

$$P_S e_{pq} P_S = e_{pq} P_S \quad (p \le q). \tag{30}$$

For the Cholesky factor R, both R and δR are upper triangular, as is $L = \delta R R^{-1}$. The coefficient direction is therefore

$$\delta Y = \mathcal{T}\mathcal{W}\delta\mathbf{C} + e(L)Y, \qquad e(L) = \sum_{p\le q} L_{pq} e_{pq}, \qquad P_S\delta Y = \delta Y. \tag{31}$$

Thus the same support S contains accepted states, selected-coefficient responses and their Cholesky representation directions. Their Hamiltonian channel sweeps need its one-generator closure I, not a second closure introduced merely by differentiating the upper-triangular transform.

Only the upper-triangular part of the residual-generator cotangent is consumed by the Cholesky overlap pullback. For any residual r, including one that does not vanish outside the selected space,

$$\Gamma_{pq} = \langle r, e_{pq}Y\rangle = \langle P_S r, e_{pq}Y\rangle, \qquad p \le q, \tag{32}$$

$$\delta\Gamma_{pq} = \langle P_S\delta r, e_{pq}Y\rangle + \langle P_S r, e_{pq}\delta Y\rangle, \qquad p \le q. \tag{33}$$

These identities permit a streamed physical overlap pullback to consume residual panels restricted to S, without retaining a complete residual carrier. They do not justify dropping the lower-triangular components from an interface that promises the full oriented generator cotangent. In general those components change under P_S and can be nonzero.

This is an invariant coordinate *support* for upper-triangular actions, not closure of an arbitrarily selected VB structure space. Projection to selected structure coefficients before a residual-generator contraction remains invalid. The argument uses neither equal alpha/beta electron counts nor a common exchange parity.

4.5 Separation of the two spin factors

Define the same-spin channel actions

$$D_\sigma = \sum_{PQ} K_{PQ} E_{\sigma P} E_{\sigma Q}. \tag{34}$$

They include the particular distribution of one-electron terms implied by Eq. (23); they are not independently assigned physical α- and β-spin Hamiltonians. Expanding the two generators gives

$$\begin{aligned} \mathcal{E}_P \mathcal{E}_Q (LF^\mathsf{T}) = (E_{\alpha P} E_{\alpha Q} L) F^\mathsf{T} + L (E_{\beta P} E_{\beta Q} F)^\mathsf{T} \\ + (E_{\alpha P} L)(E_{\beta Q} F)^\mathsf{T} + (E_{\alpha Q} L)(E_{\beta P} F)^\mathsf{T}. \end{aligned} \tag{35}$$

The last two terms become identical after summation because $K_{PQ} = K_{QP}$. Consequently,

$$\boxed{\mathcal{H}(LF^\mathsf{T}) = (D_\alpha L) F^\mathsf{T} + L (D_\beta F)^\mathsf{T} + 2 \sum_{PQ} K_{PQ} (E_{\alpha P} L)(E_{\beta Q} F)^\mathsf{T}.} \tag{36}$$

The proof uses commutativity of different spin sectors, not commutativity of generators within one spin sector. It applies to unequal U_α, U_β, unequal electron counts, ionic structures, and arbitrary mixtures of selected structures.

5 Direct sampling and structure projection

Introduce the pulled-back factors

$$L_0 = T_\alpha^\mathsf{T} L, \qquad F_0 = T_\beta^\mathsf{T} F, \tag{37}$$

$$L_H = T_\alpha^\mathsf{T} D_\alpha L, \qquad F_H = T_\beta^\mathsf{T} D_\beta F, \tag{38}$$

$$L_P = T_\alpha^\mathsf{T} E_{\alpha P} L, \qquad G_P = T_\beta^\mathsf{T} \sum_Q K_{PQ} E_{\beta Q} F. \tag{39}$$

Combining Eqs. (8) and (36), the needed original product entries are

$$m_{ab} = \sum_{\ell=1}^{r} (L_0)_{a\ell} (F_0)_{b\ell}, \tag{40}$$

$$h_{ab} = \sum_{\ell=1}^{r} \left[(L_H)_{a\ell} (F_0)_{b\ell} + (L_0)_{a\ell} (F_H)_{b\ell} + 2 \sum_P (L_P)_{a\ell} (G_P)_{b\ell} \right], \qquad (a,b) \in \Omega. \tag{41}$$

For any Davidson input $\mathbf{v}$, build $X = \mathcal{W}\mathbf{v}$ and return

$$(M\mathbf{v})_I = \sum_{(a,b)\in\Omega} w_{ab,I} m_{ab}, \qquad (H\mathbf{v})_I = \sum_{(a,b)\in\Omega} w_{ab,I} h_{ab}. \tag{42}$$

Proposition 2 (Exact factored selected-space action)**.** *Equations* (15)–(42) *evaluate Eq.* (9) *exactly for arbitrary selected structures, without constructing a complete transformed product array.*

Proof. Equation (15) is the exact decomposition into nonzero original columns. Linearity gives Eq. (16), and expansion of the Hamiltonian gives Eq. (36). Applying $\mathcal{T}^*$ to each outer product yields Eqs. (40) and (41). All other final entries have zero coefficient in $\mathcal{W}^*$ by Eq. (5). The final sampling therefore omits no contribution to Eq. (9). □

The contraction can be read as the following sequence:

$$\mathbf{v} \xrightarrow{\mathcal{W}} X = AB^{\mathsf{T}} \xrightarrow{T_\alpha, T_\beta} LF^{\mathsf{T}} \xrightarrow{\text{spin generators}} \text{factor images} \xrightarrow{T_\alpha^{\mathsf{T}}, T_\beta^{\mathsf{T}}} (h, m)|_\Omega \xrightarrow{\mathcal{W}^*} (H\mathbf{v}, M\mathbf{v}). \quad (43)$$

The middle objects are represented by factors and channel panels, not by a selected-space approximation to the transformed wave function. This is how Slater–Condon evaluation remains available without requiring selected structures to form an invariant space.

6 Two additional exact reductions

6.1 Common spin-exchange parity

An optional reduction is available when the two one-spin bases and their orbital maps coincide, and every selected structure has the same certified transpose parity $\eta = \pm 1$. Define

$$\Pi X = X^{\mathsf{T}}, \qquad P_\eta = \tfrac{1}{2}(\mathcal{I} + \eta\Pi), \qquad \Pi\mathcal{W} = \eta\mathcal{W}. \quad (44)$$

Construct X_{rep} by keeping $a < b$ with coefficient $2X_{ab}$, keeping diagonal entries with coefficient X_{aa} for $\eta = +1$, and setting the other entries to zero. Odd parity has zero diagonal. Then

$$P_\eta X_{\text{rep}} = X. \quad (45)$$

Since $\mathcal{T}$ and a spin-independent $\mathcal{H}$ commute with exchange, their original-basis composite $\mathcal{A} = \mathcal{T}^*\mathcal{H}\mathcal{T}$ also commutes. Furthermore $\mathcal{W}^* P_\eta = \mathcal{W}^*$. It follows that

$$\mathcal{W}^*\mathcal{A}X_{\text{rep}} = \mathcal{W}^* P_\eta \mathcal{A}X_{\text{rep}} = \mathcal{W}^*\mathcal{A}P_\eta X_{\text{rep}} = \mathcal{W}^*\mathcal{A}X. \quad (46)$$

The same proof applies to $\mathcal{T}^*\mathcal{T}$ for the overlap. Source factorization may therefore use X_{rep}, but the final output sampling still uses *both* halves of Ω. The unprojected image of X_{rep} is not itself a parity eigenvector.

For a single structure with m independent covalent pairs, this can reduce the source count from 2^m to 2^{m-1}. It does not imply that the touched column count always halves for an arbitrary selected space. Equal spin dimensions alone are insufficient; common parity must be checked from the structure map. Parity is not a substitute for total-spin adaptation. Without it, the unreduced factorization remains exact.

6.2 Causal rows of the exterior adjoint

Only A_Ω rows of the α pullbacks and J rows of the β pullbacks enter Eq. (40). An elementary forward exterior shear updates $z_j \leftarrow z_j + s\epsilon z_i$, with orbital coefficient s and fermionic sign $\epsilon = \pm 1$. Its adjoint updates $z_i \leftarrow z_i + s\epsilon z_j$. Starting from the requested final rows, traverse the adjoint operations in reverse chronological order. Retain an update exactly when its destination is needed, and then mark its source as needed. Execute the retained updates in their original adjoint order, including the diagonal exterior scaling.

For a row selector Q_A, this computes $Q_A T_\sigma^{\mathsf{T}} Z$ exactly. Every intermediate row that can influence the answer is retained. The pruning concerns paths through a linear transformation, not closure of the selected VB space and not the magnitude of an amplitude. A bitmap on one-spin rows and retained shear-edge indices are sufficient; no product array or dense compound matrix is required.

7 Tiling, storage, and arithmetic

7.1 Streaming the exact equations

Split the r source columns into panels of width s, and the d channels into panels of width p. Different source columns are summed independently in Eq. (40); the Hamiltonian action never couples source column labels to one another. For each source panel:

1. Build its sparse A, B, and apply T_α, T_β to obtain L, F.
2. For a channel panel $\mathcal{P}$, scatter $E_{\sigma P}$ on the corresponding factors.
3. Stream channel panels Q to accumulate $V_{\sigma P} = \sum_Q K_{PQ} E_{\sigma Q} F_\sigma$, where $F_\alpha = L$ and $F_\beta = F$. This is a channel GEMM on flattened spin-row/source-column panels.
4. Gather $\sum_{P\in\mathcal{P}} E_{\sigma P} V_{\sigma P}$ into the same-spin images. Pull back $E_{\alpha P} L$ and $V_{\beta P}$ and add the cross-spin dot products on Ω.
5. After all channel panels, pull back the base and same-spin factors, add their samples, and reuse the slabs for the next source panel.

Finally apply $\mathcal{W}^*$ once to the accumulated support arrays. Source, channel, and right-hand-side loops reuse shared backing allocations; threads do not receive private copies of the carrier.

The explicit source/channel numerical slabs require

$$O\big((U_\alpha + U_\beta) s p\big) \tag{47}$$

storage, plus $O(t)$ support coefficients and images per processed vector. Additional persistent storage includes the sparse structure map, one-spin link/exterior metadata, and the $d \times d$ kernel K. The source slabs do not include any $U_\alpha U_\beta$ coefficient array. This is a workspace statement, not a bound on the total process memory. For an overlap-only action no channel panels or same-spin Hamiltonian images are required, and the factor workspace reduces to $O((U_\alpha + U_\beta)s)$.

7.2 Where a speed advantage is obtained

Let q_σ count elementary exterior shear edges. With a full channel panel ($p = d$), a leading-work bound per input vector is

$$\begin{aligned} O\Big(& r(U_\alpha + U_\beta) d^2 + r d (q_\alpha + q_\beta) + t r d \\ & + r \sum_\sigma U_\sigma n_\sigma (n - n_\sigma + 1) + \mathrm{nnz}(\mathcal{W}) \Big). \end{aligned} \tag{48}$$

The terms describe channel mixing, exterior pullbacks, final sampling, Slater–Condon spin links, and the initial/final structure contractions. The displayed shear count is an upper bound; causal pruning reduces the selected-row pullbacks. Forward base-factor transforms remain necessary. For $p < d$, channel mixing keeps the same leading arithmetic order, while spin scatters can repeat across channel panels by a factor up to $\lceil d/p \rceil$. A smaller workspace does not change any equation or variational space.

A generated-product contraction instead has channel-mixing work of order $N_{\mathrm{cl}} d^2$, where N_{cl} counts the generated product intermediates needed by that contraction order. When this support is broad and $r \ll \min(U_\alpha, U_\beta)$, the factorized order replaces the product dimension by a source-weighted sum of one-spin dimensions. For $U_\alpha \simeq U_\beta = U$ and $N_{\mathrm{cl}} \simeq U^2$, the leading mixing ratio is approximately $U/(2r)$, before other work is counted.

For broad source support, r can approach U and the factorized order need not be faster. Both orders evaluate the same orthogonal-CI action. An arithmetic planner may select between them without a memory-based admission test for the physical selected space. Exact applicability to arbitrary structures therefore does not mean that their costs are equal or that selecting one structure universally yields rank one. There is no claim of eliminating the exponential growth of the complete one-spin spaces.

8 Integral adjoints without a product carrier

Let the bra and ket amplitudes have exact factorizations

$$Y_\ell = L_\ell F_\ell^{\mathsf{T}}, \qquad Y_r = L_r F_r^{\mathsf{T}}, \qquad \mathscr{E} = \langle Y_\ell, \mathcal{H} Y_r \rangle_{\mathrm{F}}\,. \tag{49}$$

Their source counts r_ℓ and r_r need not be equal. At fixed factors, the derivative with respect to the Hamiltonian channel kernel is

$$\frac{\partial \mathscr{E}}{\partial K_{PQ}} = \left\langle \mathcal{E}_P Y_\ell, \mathcal{E}_Q Y_r \right\rangle_{\mathrm{F}} = G_{PQ}. \tag{50}$$

This uses the self-adjoint symmetric generators, but does not assume that the bra and ket, their source counts, or their spin spaces coincide.

First contract the unacted ket factors into the bra amplitude:

$$W_\alpha = Y_\ell F_r = L_\ell (F_\ell^{\mathsf{T}} F_r), \qquad W_\beta = Y_\ell^{\mathsf{T}} L_r = F_\ell (L_\ell^{\mathsf{T}} L_r). \tag{51}$$

For each channel P, form the two channel covectors

$$C_{\alpha P} = E_{\alpha P} W_\alpha + L_\ell \big(F_\ell^{\mathsf{T}} (E_{\beta P} F_r)\big), \tag{52}$$

$$C_{\beta P} = E_{\beta P} W_\beta + F_\ell \big(L_\ell^{\mathsf{T}} (E_{\alpha P} L_r)\big). \tag{53}$$

Their dimensions are $U_\alpha \times r_r$ and $U_\beta \times r_r$. Self-adjointness moves the crossed generator from the bra to the ket; for example, $(E_{\beta P} F_\ell)^{\mathsf{T}} F_r = F_\ell^{\mathsf{T}} (E_{\beta P} F_r)$. Thus no channel scatter on the bra source family is required, even when it contains many panels. The full kernel derivative is then

$$\boxed{G_{PQ} = \left\langle C_{\alpha P}, E_{\alpha Q} L_r \right\rangle_{\mathrm{F}} + \left\langle C_{\beta P}, E_{\beta Q} F_r \right\rangle_{\mathrm{F}}\,.} \tag{54}$$

To verify this identity, expand $\mathcal{E}_Q Y_r$ as $(E_{\alpha Q} L_r) F_r^{\mathsf{T}} + L_r (E_{\beta Q} F_r)^{\mathsf{T}}$ and move each unacted ket factor to the bra side of the Frobenius product. The resulting bra contractions are precisely Eq. (52). This is the reverse contraction of the same one-body channel graph used for sigma; no product-space one-body response or higher-order density tensor is formed.

Only $\operatorname{sym} G = (G + G^{\mathsf{T}})/2$ contributes to symmetric integral variations. Using the packed diagonal selector d_P and the effective one-electron matrix f defined above, Eq. (23) gives

$$\delta \mathscr{E} = \tfrac{1}{2} \langle \operatorname{sym} G, \delta V \rangle_{\mathrm{F}} + \sum_P \frac{[(\operatorname{sym} G) d]_P}{N}\, \delta f_P, \qquad \delta f_{ps} = \delta h_{ps} - \tfrac{1}{2} \sum_q \delta(pq|qs). \tag{55}$$

Unpacking the symmetric pair entries with half weight on each off-diagonal matrix element gives the full-matrix Frobenius gradients for h and V. The product and factor algorithms share this final kernel pullback in code; they differ only in the contraction order that generates G.

8.1 A bounded reverse sweep

For each ket source panel, stream the bra factors to accumulate W_α, W_β once. For each P-channel panel, scatter the ket and these weighted factors, then stream the bra base factors to accumulate the crossed contributions to $C_{\sigma P}$. Only after the complete bra has been consumed, stream the ket Q channels and add their Gram matrices to G. Putting the Q contraction inside the bra source loop would repeat the leading channel GEMM unnecessarily; that order is not used.

Let s_ℓ, s_r be the bra and ket panel widths and p the channel width. The required numerical workspace, measured in doubles, is bounded by

$$(U_\alpha + U_\beta)\big[s_\ell + 2 s_r + (2 + \mathbf{1}_{p<d}) s_r p\big] + s_\ell s_r p + \max(\omega_\ell s_\ell, \omega_r s_r). \tag{56}$$

It includes loaded bra/ket base factors, weighted bra factors, ket channel panels, channel covectors, and one reusable source moment matrix. When $p = d$, the ket P panel also supplies all Q channels; otherwise a separate Q panel is required. Here ω_ℓ, ω_r count temporary doubles per source column used by the respective factor loaders, excluding their output views. They vanish for the original exterior-factor provider; source generation is sequential, so the larger loading scratch, not their sum, enters the bound. The immutable metadata and the required returned $d \times d$ gradient are separate. Both source and channel widths can decrease to respect the budget, without changing the represented amplitudes or a derivative tolerance. Source and channel widths are chosen jointly by the repeated-work estimate in Section 8.8. A complete bra is retained when that schedule makes it worthwhile; otherwise the same equations stream its panels.

With complete channel panels, the leading dense arithmetic is

$$O\left((U_\alpha + U_\beta)[r_r d^2 + r_\ell r_r d]\right), \tag{57}$$

instead of product-carrier Gram work of order $U_\alpha U_\beta d^2$. The source-pair term is essential and is not assumed linear in rank. Self-adjointness permits exchanging the bra and ket before contraction, placing the smaller source family on the ket side. One-spin link traversals and panel regeneration also contribute; a smaller channel panel can repeat these traversals. No universal speedup is inferred from the leading dense terms alone.

For an exchange-parity reduction, a representative cannot replace both amplitudes without projection. A physical bra with parity η may, however, contract an unprojected ket representative Z when $Y_r = P_\eta Z$ and the integral derivative commutes with exchange: $\langle Y_\ell, \delta\mathcal{H} Z\rangle_{\mathrm{F}} = \langle Y_\ell, \delta\mathcal{H} Y_r\rangle_{\mathrm{F}}$. The caller must establish this identity; the common factor kernel does not infer a parity projection from equal spin dimensions.

8.2 Directional factors and the integral part of an HVP

Let R be the nonsingular upper-triangular orbital map in the exterior transforms and let δR be its direction. Define the spatial generator and its one-spin representation by

$$\kappa = \delta R R^{-1}, \qquad \mathcal{G}_\sigma(\kappa) = \sum_{pq} \kappa_{pq} a^\dagger_{p\sigma} a_{q\sigma}. \tag{58}$$

The derivative of an exterior representation obeys $\delta T_\sigma = \mathcal{G}_\sigma(\kappa) T_\sigma$. This generator differentiates the CI representation; it does not replace the physical nonorthogonal VB orbital coordinates or their finite update. For $Y = LF^{\mathsf{T}}$, with selected structure coefficients allowed to vary,

$$\boxed{\delta Y = (\mathcal{G}_\alpha(\kappa)L)F^{\mathsf{T}} + L(\mathcal{G}_\beta(\kappa)F)^{\mathsf{T}} + T_\alpha\, \delta X\, T_\beta^{\mathsf{T}}, \qquad \delta X = \mathcal{W}\, \delta\mathbf{C}.} \tag{59}$$

The structure map $\mathcal{W}$ and its exact support plan are fixed. The last term therefore uses the same support factorization with the structure-response coefficients; it is not a projection of either orbital derivative onto the original support. The two orbital terms generally leave that support and retain every contribution through their transformed spin factors.

The response interface therefore carries $\delta\mathbf{C}$ in the selected structure basis, not δX as a full spin-product array. For state s and an original touched product e, it loads only

$$(\delta x_s)_e = \sum_I W_{eI}\, \delta C_{Is}. \tag{60}$$

The incidences W_{eI} and source cover are already retained at the accepted point. This costs $O(\mathrm{nnz}\, W)$ additions and scalar products per state, with $O(P)$ temporary coefficients for P touched products, instead of an $O(U_\alpha U_\beta)$ response array per state. The response columns are borrowed from their owning structure-space calculation; they are not copied into a second bundle. Both the integral-gradient derivative and the

oriented residual-generator derivative consume these same columns. The optional Hamiltonian/overlap output projection is omitted when only the response adjoint is requested. This changes neither the full orbital derivative in Eq. (59) nor the selected-space variational problem. In particular, a zero-dimensional horizontal CI tangent does not remove the normalization response or its gauge multiplier.

For each coordinate-anchored source j, write $Y_j = \ell_j f_j^\mathsf{T}$. Its selected-coefficient derivative changes only the unanchored raw factor. For a β-anchored source $\delta f_j = 0$; for an α-anchored source $\delta \ell_j = 0$. The symbols $\delta \ell_j$ and δf_j here denote coefficient changes at fixed orbital map. Combining these changes with the orbital derivative gives

$$a_j = \mathcal{G}_\alpha(\kappa)\ell_j + \delta \ell_j, \qquad b_j = \mathcal{G}_\beta(\kappa) f_j + \delta f_j, \tag{61}$$

$$\delta Y = \sum_{j=1}^{r} \left(a_j f_j^\mathsf{T} + \ell_j b_j^\mathsf{T} \right). \tag{62}$$

Thus at most $2r$ factors are required, rather than the $3r$ produced by concatenating two orbital families and a separate coefficient-response family. No cross term $\delta \ell_j \delta f_j^\mathsf{T}$ enters a first derivative. This is an exact source identity, not a numerical rank truncation. For a physical exchange-parity projection, apply the same linear projection to both sides of Eq. (62), including exchanged anchors and their signs. A coefficient-only tangent needs only r factors.

A source panel loads its base factors and applies the oriented Slater–Condon generator to the required spin factor only. The same temporary is then overwritten by the coefficient derivative of that factor, transformed and accumulated. No second factor panel is needed. Unlike a Hamiltonian channel, κ need not be symmetric: replacing it by its symmetric part would give the wrong orbital derivative. The one-factor temporary uses at most $\max(U_\alpha, U_\beta)s$ doubles for panel width s; it is included in the loader workspace term of the preceding bound.

Write $\mathcal{J}(Y_\ell, Y_r)$ for the symmetric integral adjoint of $\langle Y_\ell, \mathcal{H} Y_r \rangle_\mathrm{F}$, including the shared K-kernel pullback. Since each physical integral variation is self-adjoint,

$$\delta \mathcal{J}(Y, Y) = 2\mathcal{J}(\delta Y, Y). \tag{63}$$

Thus the integral-gradient derivative needed in an orbital HVP can be contracted from the directional factor source directly, without constructing δY or any higher-order RDM for this contraction. If the accepted and directional factor counts are r and r_δ, the leading dense work is

$$O\left((U_\alpha + U_\beta)[\min(r, r_\delta) d^2 + r\, r_\delta d] \right), \tag{64}$$

plus the one-spin link traversals and source transforms. The count r_δ includes both orbital terms and any structure response, including the physical parity reconstruction when required. It must not be estimated from the accepted source count alone.

There are two exact placements of a certified exchange-parity projector P_η. If Z is the accepted representative and δZ its tangent, self-adjointness and $[P_\eta, \delta\mathcal{H}] = 0$ give

$$\left\langle P_\eta \delta Z, \delta\mathcal{H} Z \right\rangle_\mathrm{F} = \left\langle P_\eta Z, \delta\mathcal{H} \delta Z \right\rangle_\mathrm{F}. \tag{65}$$

For an accepted representative with r factors and a merged tangent with $2r$ factors, the two source-count pairs are $(4r, r)$ and $(2r, 2r)$. Both have the same source-pair moment count. The first has less leading channel-Gram arithmetic, but can repeatedly regenerate its larger directional bra when it does not fit the common workspace. The second can keep the accepted bra resident without enlarging that workspace. Compare the full scalar-work estimate, including the actual planned provider loads, channel traversals and Gram work, rather than selecting solely by the smaller ket. The projector placement changes only contraction order, never the selected VB space or an accuracy criterion. Equation (65) applies to the symmetric integral adjoint; it does not permit moving a Hamiltonian through the oriented generator in a residual adjoint.

Equation (63) is the integral part of the HVP, not its complete orbital pullback. The derivative of the exterior-map adjoint, the full carrier residual, and the state-average gauge terms remain required. The present implementation keeps those terms while migrating the integral contraction to Eq. (63).

8.3 Oriented generator and residual contractions

Define the orbital matrix $\mathcal{D}(X,Y)$ by

$$[\mathcal{D}(X,Y)]_{pq} = \langle X, (a^\dagger_{p\alpha}a_{q\alpha} + a^\dagger_{p\beta}a_{q\beta})Y\rangle_{\mathrm{F}}. \tag{66}$$

It is the coefficient of κ_{pq} in $\langle X, \mathcal{G}(\kappa)Y\rangle_{\mathrm{F}}$ and need not be symmetric. For $Y = L_r F_r^{\mathsf{T}}$, form only the weighted single-spin panels

$$W_\alpha = XF_r, \qquad W_\beta = X^{\mathsf{T}}L_r. \tag{67}$$

Then

$$[\mathcal{D}(X,Y)]_{pq} = \sum_j \langle (W_\alpha)_j, a^\dagger_{p\alpha}a_{q\alpha}(L_r)_j\rangle + \sum_j \langle (W_\beta)_j, a^\dagger_{p\beta}a_{q\beta}(F_r)_j\rangle. \tag{68}$$

If $X = L_\ell F_\ell^{\mathsf{T}}$, each weighted panel is obtained from source-pair moments, $XF_r = L_\ell(F_\ell^{\mathsf{T}}F_r)$ and $X^{\mathsf{T}}L_r = F_\ell(L_\ell^{\mathsf{T}}L_r)$. The final contractions traverse the oriented single-spin occupation/replacement links. They are not symmetric Hamiltonian-channel derivatives. For real coefficients, $\mathcal{D}(X,Y) = \mathcal{D}(Y,X)^{\mathsf{T}}$; this permits a transposed contraction order but not an unqualified interchange of sources.

For the residual $Z = (\mathcal{H} - E)X$, the required panels can also be obtained without Z or $\mathcal{H}X$. Using the same-spin actions D_σ and the ket channel mixtures

$$V_{\alpha P} = \sum_Q K_{PQ}E_{\alpha Q}L_r, \qquad V_{\beta P} = \sum_Q K_{PQ}E_{\beta Q}F_r, \tag{69}$$

the exact residual-weighted panels are

$$Z_\alpha = D_\alpha W_\alpha + X(D_\beta F_r) + 2\sum_P E_{\alpha P}(XV_{\beta P}) - EW_\alpha, \tag{70}$$

$$Z_\beta = D_\beta W_\beta + X^{\mathsf{T}}(D_\alpha L_r) + 2\sum_P E_{\beta P}(X^{\mathsf{T}}V_{\alpha P}) - EW_\beta. \tag{71}$$

Thus $Z_\alpha = ZF_r$ and $Z_\beta = Z^{\mathsf{T}}L_r$. Replacing W_α, W_β by Z_α, Z_β in Eq. (68) gives $\mathcal{D}((\mathcal{H}-E)X, Y)$. The spectator actions have been moved to the ket using self-adjointness of D_σ; the crossed term follows from Eq. (36). Self-adjointness of $\mathcal{H}$ does *not* justify exchanging the two sources here, because $\mathcal{H}$ generally does not commute with an oriented one-body generator.

The implementation reuses the factor/channel sweep twice: once for the weighted panels to produce $D_\sigma W_\sigma$, and once for the ket panels to produce the spectator and crossed terms. After same-spin gather, the ket's raw channel buffers are reused for $XV_{\beta P}$ and $X^{\mathsf{T}}V_{\alpha P}$, then consumed by the crossed gather. No separate channel scatter/mixing implementation or product-space residual is required.

With complete channel panels, the leading dense arithmetic is $O((U_\alpha + U_\beta)[r_r d^2 + r_\ell r_r d])$, plus the oriented single-spin link contractions. Without $\mathcal{H}$, it reduces to

$$O\left((U_\alpha + U_\beta)r_\ell r_r + r_r \sum_\sigma U_\sigma n_\sigma(n - n_\sigma + 1)\right). \tag{72}$$

Accepted-state transition densities use this estimate to choose between factor moments and product-space links. Their planner does not reuse the Hamiltonian-channel cost estimate.

A conservative numerical workspace bound for the residual contraction, in doubles, is

$$(U_\alpha + U_\beta)[s_\ell + (3p+9)s_r] + s_\ell s_r \max(1,p) + \max(\omega_\ell s_\ell, \omega_r s_r) + (t+1)n^2, \tag{73}$$

where t bounds the number of workers and the last term covers the small orbital-matrix reduction and its temporary result. The required returned $n\times n$ matrix and immutable metadata are separate. For a generator without $\mathcal{H}$, the channel sweep and residual outputs are absent, giving $(U_\alpha + U_\beta)(s_\ell + 2s_r) + s_\ell s_r$ plus the same loading and reduction terms. No per-worker spin/product array is allocated.

8.4 Reuse of the normalization response

Let C contain the selected structure coefficients and M be the structure-overlap matrix. The solve-free normalization response has the known internal state frame

$$M_C = C^\mathsf{T} M C, \qquad K_g = -\tfrac{1}{2} C^\mathsf{T} \delta M C, \qquad A_g = M_C^{-1} K_g, \qquad Q_g = C A_g. \tag{74}$$

Retain A_g from the small selected-frame solve rather than infer it from the expanded numerical vectors Q_g. This identity requires no closure of the selected VB space under an orbital transformation.

For a single selected root,

$$a = -\frac{1}{2}\frac{C^\mathsf{T}\delta M C}{C^\mathsf{T} M C}, \qquad Q_g = aC. \tag{75}$$

At fixed accepted active-space eigenvalue ϵ, the integral adjoint and the residual-generator adjoint are quadratic in C. Therefore their coefficient-only directional derivatives are

$$\delta_C \mathcal{J}(Y,Y) = 2a\, \mathcal{J}(Y,Y), \tag{76}$$

$$\delta_C \mathcal{D}((\mathcal{H}-\epsilon)Y,Y) = 2a\, \mathcal{D}((\mathcal{H}-\epsilon)Y,Y). \tag{77}$$

These identities do not assume that the full-carrier Ritz residual vanishes. Since the accepted orbital-map pullback is linear in these adjoints, their coefficient response is also $2a$ times the already retained active-integral adjoint. No new spin-source generation, Hamiltonian action or integral contraction is required for this part of the response. Given the already computed frame a, the reuse costs $O(d^2+n^2)$ scalar work in the current active-integral representation, with no extra array beyond the existing directional adjoint; n is the active orbital count and $d = n(n+1)/2$. This removes the carrier dimensions and source rank from the coefficient-scaling update, not from the rest of the orbital HVP.

The selected-state multiplier contribution remains separate, as does the directional derivative of the orbital-map pullback. Neither is removed by Eq. (77). The reused quantity is the accepted *active-integral* adjoint, not the total orbital gradient, which also contains inactive-only contributions. For multiple states, retain the full matrix A_g; the weighted diagonal adjoint totals alone do not contain the required cross-state adjoints. The general coefficient-response contraction remains necessary until those transition adjoints are available.

8.5 Selected-structure response without a response carrier

At a fixed orbital point let $Y_i = \mathcal{T}\mathcal{W}\mathbf{C}_i$ and $Q_i = \mathcal{T}\mathcal{W}\mathbf{q}_i$, where $\mathbf{q}_i$ is a coefficient response in the original selected structure space. Write $\mathcal{J}_i = \mathcal{J}(Y_i,Y_i)$ and $\Gamma_i = 2\mathcal{D}((\mathcal{H}-E_i)Y_i,Y_i)$ for its unweighted accepted integral and exterior-generator adjoints. With E_i fixed, their coefficient-response contributions are exactly

$$\delta\mathcal{J}_i = 2\mathcal{J}(Q_i,Y_i), \tag{78}$$

$$\delta\Gamma_i = 2\mathcal{D}((\mathcal{H}-E_i)Q_i,Y_i) + 2\mathcal{D}((\mathcal{H}-E_i)Y_i,Q_i). \tag{79}$$

The first residual is the response of the wave function; the second is the response of the generator's right argument. Neither residual may be projected into the selected structure space before its generator contraction. Both are evaluated by the residual panels of Eq. (70). Thus no Q_i, $\mathcal{H}Q_i$, or response residual product array is required. Each raw response is generated once on Ω and reused by the spin-factor adjoints and the tiled $\mathcal{W}^*\mathcal{T}^*\mathcal{H}Q_i$, $\mathcal{W}^*\mathcal{T}^*Q_i$ projections.

8.6 Product-tile response with resident accepted data

For a broad source graph the generated product order can be cheaper than factor-pair moments. Let $R_i = (\mathcal{H} - E_i)Y_i$ be the retained accepted residual. The transpose identity for real one-body generators is

$$\mathcal{D}(L, R) = \mathcal{D}(R, L)^{\mathsf{T}}, \qquad (e_{pq})^{\mathsf{T}} = e_{qp}. \tag{80}$$

It rewrites the coefficient-response generator cotangent as

$$\delta\Gamma_i = 2\big[\mathcal{D}(\mathcal{H}Q_i, Y_i) - E_i\mathcal{D}(Q_i, Y_i) + \mathcal{D}(Q_i, R_i)^{\mathsf{T}}\big]. \tag{81}$$

Every contraction now has a generated response panel as its left argument and a resident accepted array as its right argument. There is no need to build a complete Q_i, $\mathcal{H}Q_i$, or response residual array, nor to recompute $\mathcal{H}Y_i$. For disjoint beta-column projectors P_t, linearity gives $\mathcal{D}(L, R) = \sum_t \mathcal{D}(P_t L, R)$. The two input-panel terms in Eq. (81) share their signed-link traversal; their second term is transposed only *after* its oriented contraction. Truncating its lower triangle before transposition would lose required upper-triangular information.

The generated Hamiltonian sweep also forms the channel Gram for $\mathcal{J}(Q_i, Y_i)$. Each emitted Hamiltonian panel is consumed both by the first term in Eq. (81) and by the exterior/structure adjoint, before its storage is reused. Cache eviction may emit partial Hamiltonian panels: linearity of both consumers makes this exact without waiting for a complete sigma column. The physical Cholesky pullback needs only the upper triangle. By Eq. (33), the first term may therefore consume the Hamiltonian image restricted to the generated support S. This restriction does not assert closure in the selected structure basis. For a common exchange parity, the same summed-spin generators commute with the exchange projector, so the projected-action representative is sufficient for this physical cotangent as well as for the structure image.

For a left panel whose beta columns belong to B_t, the two spin contributions to an oriented generator cotangent can be contracted as

$$\mathcal{D}^{\alpha}_{pq}(L, R) = \sum_{aa'} (e^{\alpha}_{pq})_{aa'} \sum_{b \in B_t} L_{ab} R_{a'b}, \tag{82}$$

$$\mathcal{D}^{\beta}_{pq}(L, R) = \sum_{b \in B_t, b'} (e^{\beta}_{pq})_{bb'} \sum_{a} L_{ab} R_{ab'}. \tag{83}$$

Each signed link is visited outside the spectator dot product. Let ℓ_σ count diagonal and single replacements per spin string. The metadata work per right-hand side changes from

$$O\big(U_\alpha |B_t| (\ell_\alpha + \ell_\beta)\big) \quad \text{to} \quad O\big(U_\alpha \ell_\alpha + |B_t| \ell_\beta\big).$$

The numerical dot-product work retains the former order; this is a reduction of repeated indexing and branching, not a lower orbital power for the entire contraction. Beta contributions use contiguous column dots and alpha contributions borrow strided row views, with no transposed carrier copy. The implementation partitions alpha rows and beta link events separately, so even one beta column can share its long spectator contractions across the existing OpenMP team. Only orbital-matrix reductions are partition-local. Upper-triangle selection is performed after the requested transpose and outside the spectator dot. Exact-zero stripes may be skipped, but no amplitude threshold or norm-based small-vector rejection is used.

Compact symbolic panels. For a beta panel B_t, let $S_t = \{(a, b) \in S : b \in B_t\}$. The retained input and partial-image panels store only $|S_t| b$ floating-point values for b simultaneous right-hand sides, together with one shared row-index pattern of size $O(|S_t| + |B_t|)$. All entries in this pattern are retained, including exact zeros and arbitrarily small values. This is symbolic support compression, not numerical sparsification. The Slater–Condon gather adds directly into these indexed rows. An ordered exterior

adjoint temporarily unpacks one bounded panel; it does not unpack the complete carrier. Full symbolic support uses implicit row indices. Both caches retain their byte limits, and partial-image eviction remains linear as above. Shared patterns are held only by resident panels, rather than by an additional unbounded support table.

For a complete covalent family with $n = 2m$ orbitals and m electrons per spin, the generated support obeys the prefix condition $N_\alpha(k) + N_\beta(k) \geq k$ for every k. Its cardinality is

$$|S_{\mathrm{cov}}| = \binom{2m}{m}^2 - \binom{2m}{m+1}^2. \tag{84}$$

To count the excluded pairs, reflect both occupation bits, $(a_j, b_j) \mapsto (1 - b_j, 1 - a_j)$, up to the first prefix with $N_\alpha(k) + N_\beta(k) = k - 1$. This maps an excluded pair bijectively to two strings with $m + 1$ electrons; the inverse reflects up to the first prefix excess of one. Any selected covalent subset has support contained in this set. Thus the complete-covalent upper bounds for 12 and 14 orbitals are 226512 and 2760615 stored coordinates, respectively, rather than 853776 and 11778624 complete-product coordinates. This bound does not apply to arbitrary ionic selections: their own generated support determines the retained count, and the same bounded-cache algorithm still applies.

Cache misses may occur inside the persistent sigma parallel region. The ordered exterior kernels therefore share its existing workers through disjoint-column or disjoint-row tasks, instead of opening a nested parallel region that executes serially. Each task borrows the same panel; there is no per-worker floating-point carrier. This changes scheduling, not the order of elementary shears acting on any coefficient.

Accepted arrays Y_i and R_i are borrowed, not copied. Removing their storage is a separate migration of accepted and orbital-direction consumers; the present identities remove the additional *response* carriers. Original-support incidence/value buffers, bounded column caches, small orbital reductions and the active-integral cotangent have separate storage costs. In particular, a touched-support buffer is not a guarantee of subquadratic memory if the original support itself fills the product grid.

For positive state weights w_i the contributions above are summed with w_i. The separate overlap-multiplier term is

$$2 \sum_{ij} w_i \mu_{ji} \mathcal{D}(Y_j, Y_i), \qquad \mu_{ji} = \frac{1}{2}\left[\lambda_{ji} + \frac{w_j}{w_i}\lambda_{ij}\right], \tag{85}$$

where λ_{ij} are the directional normalization/orthogonality multipliers in the structure-response equations. This weight-symmetrized gauge contribution is not part of either residual panel. An exactly zero $\mathbf{q}_i$ contributes no images or derivative tensors; this follows from linearity, not an amplitude threshold.

For source counts r_ℓ, r_r, the leading dense work of the integral adjoint and both residual-generator contractions is

$$O\big((U_\alpha + U_\beta)[r_r d^2 + r_\ell r_r d]\big). \tag{86}$$

The two Hamiltonian images also supply their generator adjoints by the shared image consumer in Section 8.7. Only $\mathcal{H}Q_i$ is projected to structure H/M images; $\mathcal{H}Y_i$ is consumed by its adjoint without an unused structure projection. With equal representative source counts $r_r = r$ and physical bra count $r_\ell = l$, the leading multiply–add count used for this consumer is

$$W_{\mathrm{source}} = 6(U_\alpha + U_\beta) r d^2 + (U_\alpha + U_\beta) l r (12d + 20). \tag{87}$$

The coefficients count one integral adjoint and two shared H-image/generator adjoint sweeps. The energy shifts use the transpose identity for the ordinary generator adjoint. The fused product sigma/Gram has leading work $4N_{\mathrm{product}} d^2$, with N_{product} reduced only by a certified exchange parity. Comparing these leading operation counts does not fit timings or change numerical thresholds. Link work, source transforms, sampling and finite tile effects still affect wall time. A product order can still be cheaper for broad support; this is an exact arithmetic-order choice, not a change of variational space or a memory admission criterion.

8.7 Accepted states and orbital directions without product arrays

Let $F = \mathcal{T}\mathcal{W}$, $Y = F\mathbf{C}$ and $\kappa = \delta R R^{-1}$. Holding the structure coefficients fixed,

$$\delta(H\mathbf{C}) = F^*\big[\delta\mathcal{H}Y + \mathcal{H}\mathcal{G}(\kappa)Y + \mathcal{G}(\kappa)^*\mathcal{H}Y\big], \tag{88}$$

$$\delta(M\mathbf{C}) = F^*\big[\mathcal{G}(\kappa) + \mathcal{G}(\kappa)^*\big]Y. \tag{89}$$

The left-generator term differentiates the projection, not the physical selected structure map. It may not be omitted even if the selected coefficients solve their generalized eigenproblem.

Each Hamiltonian image is a sum of outer products from Eq. (36). For any such term AB^{T},

$$\mathcal{G}(\kappa)^*(AB^{\mathsf{T}}) = [\mathcal{G}_\alpha(\kappa^{\mathsf{T}})A]B^{\mathsf{T}} + A[\mathcal{G}_\beta(\kappa^{\mathsf{T}})B]^{\mathsf{T}}. \tag{90}$$

Apply the oriented spin generators to the factor panels, perform the same causal exterior adjoints, and sample only Ω. Two additional shared spin-factor channel slabs suffice, including both base and same-spin images. They are reserved inside the original numerical budget, never allocated independently per worker.

For a general orbital/structure tangent δY, the accepted generator adjoint $\Gamma = 2\mathcal{D}((\mathcal{H} - E)Y, Y)$ has direction

$$\delta\Gamma = 2\mathcal{D}((\mathcal{H} - E)\delta Y, Y) + 2\mathcal{D}((\mathcal{H} - E)Y, \delta Y) + 2\mathcal{D}(\delta\mathcal{H}Y, Y), \tag{91}$$

before adding the separate normalization/orthogonality multipliers. Here E is held fixed; its required directional multiplier terms follow the same state-average gauge convention as Section 8.5. The third term uses the residual-panel kernel with Hamiltonian $\delta\mathcal{H}$ and shift zero. The integral direction remains $2\mathcal{J}(\delta Y, Y)$. Consequently the accepted adjoints and their complete directions require only structure coefficients, source providers, small orbital tensors and bounded numerical panels, not accepted amplitudes or residual product arrays. The differentiated Cholesky/integral pullback is still applied afterwards; Eq. (91) does not replace that chain rule.

Computing Eq. (91) after releasing the directional H/M images would repeat their Hamiltonian work. Instead consume each outer-product image $X = \sum_j A_j B_j^{\mathsf{T}}$ before its exterior projection. For right factors $Y_r = L_r F_r^{\mathsf{T}}$, accumulate

$$Z_\alpha \mathrel{+}= A(B^{\mathsf{T}}F_r), \qquad Z_\beta \mathrel{+}= B(A^{\mathsf{T}}L_r). \tag{92}$$

After all source/channel tiles, these are XF_r and $X^{\mathsf{T}}L_r$; Eq. (68) then returns $\mathcal{D}(X, Y_r)$. For a certified exchange representative, accumulate its physical image $P_\eta X = \frac{1}{2}(X + \eta X^{\mathsf{T}})$ by adding the exchanged factors with weight $\eta/2$. No product-space symmetrization is formed.

Let $D = \mathcal{D}(Y, \delta Y)$. Since $\mathcal{D}(\delta Y, Y) = D^{\mathsf{T}}$, the orbital-direction cache is only

$$\delta\Gamma = 2\big[\mathcal{D}(\mathcal{H}\delta Y, Y) + \mathcal{D}(\delta\mathcal{H}Y, Y) + \mathcal{D}(\mathcal{H}Y, \delta Y) - E(D + D^{\mathsf{T}})\big]. \tag{93}$$

The three Hamiltonian images are already generated for Eq. (88). Their generator adjoints are collected in that sweep; the subsequent HVP pullback reuses this small $n \times n$ covector, not a sigma or residual product array. With state averaging, the cache is the weighted sum and coefficient-response/gauge terms are added separately. The integral-adjoint direction remains required.

One numerical budget is split between the H-image workspace and the right factor/moment workspace. Right factors are tiled if necessary; the H source can then replay, but its structure projection occurs only on the first pass. When all right factors fit, the three images visit $4r$ representative source columns: $2r$ for $\mathcal{H}\delta Y$, r for $\delta\mathcal{H}Y$, and r for $\mathcal{H}Y$. The unfused residual-panel route plus the images visits $12r$. This removes repeated channel mixing, not the source-pair moment arithmetic or the outer optimizer's HVP calls.

8.8 Joint source/channel scheduling

For source widths s_ℓ, s_r and channel width p, define

$$N_\ell = \lceil r_\ell / s_\ell \rceil, \qquad N_r = \lceil r_r / s_r \rceil, \qquad N_P = \lceil d/p \rceil, \qquad U = U_\alpha + U_\beta. \tag{94}$$

The leading dense channel and moment arithmetic does not depend on these widths, but source regeneration, link traversal and output traffic do. In particular, the number of loaded bra source columns is

$$C_\ell = \begin{cases} r_\ell, & s_\ell = r_\ell, \\ N_r(N_P + b)r_\ell, & s_\ell < r_\ell, \end{cases} \qquad b = \begin{cases} 1, & \text{integral adjoint}, \\ 2, & \text{residual generator}. \end{cases} \tag{95}$$

Here b counts the channel-independent moment sweeps. If generating a source column costs ℓ_ℓ nominal scalar operations and element touches, source work is $C_\ell \ell_\ell$. Each provider declares this work from its sparse writes, diagonal scales and exterior shear multiply–adds; directional providers add the oriented generator work. This metadata affects only the evaluation schedule, never the amplitudes, ranks or tolerances.

The single-replacement links are generated in ascending packed-channel order. For each larger orbital index, ascending lower indices with opposite occupation generate every link once. A channel tile then uses binary ranges in the existing link and occupation arrays. It does not scan all links and reject those outside the tile, nor allocate a string-by-channel offset table. Let $J = \sum_\sigma U_\sigma n_\sigma (n - n_\sigma + 1)$ count the single-spin occupation/replacement links. Across all source columns, matching-link visits scale as $r_r v(N_P) J$, where

$$v_I(N_P) = N_P + 1, \qquad v_R(N_P) = 2N_P + 3. \tag{96}$$

Without channel ranges these counts were respectively $N_P^2 + N_P$ and $2N_P^2 + 3N_P$. Binary lookup comparisons still repeat with those latter counts and are included through an upper comparison estimate B per complete single-spin row family. The complete-channel case needs no binary lookup penalty.

The schedule also counts repeated output reads and writes in the moment GEMMs, and repeated channel zero fills:

$$W_{\text{out}} = 2U r_r (d + b) N_\ell, \tag{97}$$

$$W_{0,I} = U r_r d (N_P + 1), \qquad W_{0,R} = 2U r_r d N_P. \tag{98}$$

Combining these terms gives the nominal variable-work estimate

$$\widehat{W} = C_\ell \ell_\ell + r_r \left[3v(N_P) J + s(N_P) B \right] + W_0 + W_{\text{out}}, \tag{99}$$

where $s_I = N_P^2 + N_P$, $s_R = 2N_P^2 + 3N_P$, and the factor three counts one matching-link visit plus its multiply and add. Terms invariant across the candidate schedules are omitted. The estimate combines scalar work and element traffic without fitted hardware weights; it is not a prediction of absolute wall time, which requires native measurement.

The planner searches feasible source and channel widths under the same memory bound. For each ket/channel choice it takes the largest feasible bra panel; ties use fewer moment panels. Partial-channel residual sweeps currently use one ket column because the shared forward workspace and moment buffers must use the same channel width. Full-channel sweeps allow wider ket panels. This is a workspace scheduling limit, not a restriction on the selected VB space.

9 A shared accepted–tangent product sweep

For broad source support, retain the generated product contraction order without retaining unbounded accepted or directional product arrays. Write $F = \mathcal{T}\mathcal{W}$, $Y = F\mathbf{C}$, and let δF denote the change of the

exterior map under the orbital direction. The complete directional wave function and the accepted residual are

$$Z = \delta F\mathbf{C} + F\delta\mathbf{C}, \qquad R_Y = (\mathcal{H} - E)Y. \tag{100}$$

Here E is held fixed in the orbital Lagrangian; its response and the state-average gauge multipliers are included separately as in Section 8.5. Differentiating the complete projection, rather than projecting an intermediate wave function, gives

$$\delta(F^*\mathcal{H}Y) = F^*(\delta\mathcal{H}Y + \mathcal{H}Z) + (\delta F)^*\mathcal{H}Y, \tag{101}$$

$$\delta(F^*Y) = F^*Z + (\delta F)^*Y. \tag{102}$$

The left-map derivative is necessary even when $\delta\mathbf{C} = 0$. No invariance of the selected structure space is used.

With $\mathcal{J}$ and $\mathcal{D}$ defined above, the accepted and directional orbital-generator cotangents are

$$\Gamma = 2\mathcal{D}(R_Y, Y), \tag{103}$$

$$\delta\mathcal{J} = 2\mathcal{J}(Z, Y), \tag{104}$$

$$\delta\Gamma = 2\left[\mathcal{D}(\mathcal{H}Z + \delta\mathcal{H}Y, Y) + \mathcal{D}(R_Y, Z) - E\mathcal{D}(Z, Y)\right]. \tag{105}$$

Only the upper triangle of the final generator cotangent is needed by the Cholesky pullback. The previously proved support certificate permits R_Y and the emitted Hamiltonian images to be restricted to S for these upper cotangents. It does not permit replacing a general oriented transition density by its upper triangle.

If compact accepted panels of (Y, R_Y) are resident, the same panels supply the right readers, $\mathcal{D}(R_Y, Z)$, and $\mathcal{H}Y = R_Y + EY$ in the left projection of Eq. (101). Thus a coefficient-only response requires neither another accepted exterior transform nor another accepted Hamiltonian action. The channel sweep producing $\mathcal{H}Z$ also accumulates $\mathcal{J}(Z, Y)$ before its slabs are released. The forward orbital-direction consumer retains the resulting small cotangents and prepared integral tangent; its reverse consumer reuses them.

When accepted panels cannot all remain resident under the shared budget, generate Y and Z together and apply $\mathcal{H}$ to this two-column-family block. The equivalent unretained form of the generator response is

$$\delta\Gamma = 2\left[\mathcal{D}(\mathcal{H}Z + \delta\mathcal{H}Y, Y) + \mathcal{D}(\mathcal{H}Y, Z) - E\{\mathcal{D}(Z, Y) + \mathcal{D}(Y, Z)\}\right]. \tag{106}$$

For real wave functions the last brace equals $\mathcal{D}(Z, Y) + \mathcal{D}(Z, Y)^{\mathsf{T}}$. One complete oriented traversal therefore supplies both terms, followed by upper-triangle selection. The resident and streamed equations are identical; residency changes only reuse, not the variational space or accuracy. This does not assert that all accepted work can be retained at fixed memory without recomputation.

Each emitted Hamiltonian panel is consumed by the structure projection and generator adjoint before release. Partial images emitted on eviction are valid because both consumers are linear. A one-body generator's neighbor columns are themselves read in bounded batches; their union must not become a complete right carrier. With panel width b, shared residency bound B and the existing bounded sigma workspace W_σ, auxiliary numerical storage, excluding immutable incidence and one-spin topology, is

$$O\left(t + (U_\alpha + U_\beta)b + B + W_\sigma + n^4\right) \tag{107}$$

per sequentially processed state, in addition to the selected coefficient block. The n^4 term denotes the retained active integral and adjoint tensors, not a spin-product array. If original touched support is complete, t itself can equal $U_\alpha U_\beta$; this bound does not conceal that input cost. Accepted residency is shared across roots, never copied per worker.

10 Necessary projections and cancelling coordinate maps

An overlap solve is not a physical observable. Its necessity depends on the coordinate map that the optimizer has chosen. Let $F(p) = \mathcal{T}(p)\mathcal{W}$ map selected structure coefficients into an orthogonal spin-string carrier at orbital displacement p. With an independent structure displacement q, write $\mathbf{C}_b = \mathbf{C}_0 + Zq$, where Z embeds the structure tangent. The current finite CI transport projects the target

$$b(p, q) = F(0)\mathbf{C}_b + \dot{F}(0)[p]\mathbf{C}_0 \tag{108}$$

onto the selected space at the displaced orbitals. Its normal equation is

$$M(p)\mathbf{C}(p, q) = F(p)^{\mathsf{T}} b(p, q), \qquad M(p) = F(p)^{\mathsf{T}} F(p). \tag{109}$$

For a nonclosed selection this is a genuine restricted projection: replacing $M(p)$ by the identity, or using an inverse on a larger complete space, does not give the same transport. Its gradient pullback involves the adjoint solve with the same $M(p)$. These operations are necessary for this finite coordinate map, not for every possible VBSCF optimization framework. Changing the map requires deriving its derivatives consistently, rather than independently deleting the projection from a keyframe or an HVP. Two requests for exactly the same (p, q) within one immutable macro model, however, solve exactly the same equation; the previously computed point is sufficient for both consumers.

There is a stronger reduction when $\operatorname{ran}\mathcal{W}$ is a complete spin sector. Define its orthogonal projector P and its orbital-independent Gram G by

$$G = \mathcal{W}^*\mathcal{W}, \qquad P = \mathcal{W}G^{-1}\mathcal{W}^*. \tag{110}$$

The common spatial orbital transformation commutes with the total-spin projector. Thus $\mathcal{T}P = P\mathcal{T}$, as do its inverse and adjoint, and the Moore–Penrose inverse of $F = \mathcal{T}\mathcal{W}$ is

$$F^+ = G^{-1}\mathcal{W}^*\mathcal{T}^{-1}, \qquad FM^{-1} = \mathcal{T}^{-*}\mathcal{W}G^{-1}. \tag{111}$$

Indeed, the proposed left inverse times F is the identity, while $FF^+ = \mathcal{T}P\mathcal{T}^{-1} = P$ is an orthogonal projector. Completeness here means completeness of the selected spin sector, not merely completeness of the ambient alpha and beta string lists. Covalent-only and arbitrary selected spaces generally do not meet this hypothesis.

Let D be the existing carrier diagonal, including its root-energy shift and the unchanged denominator protection. The symmetric structure-space inverse preconditioner can therefore be evaluated as

$$M^{-1}F^*D^{-1}FM^{-1} = G^{-1}\mathcal{W}^*\mathcal{T}^{-1}D^{-1}\mathcal{T}^{-*}\mathcal{W}G^{-1}. \tag{112}$$

No commutation between D and the spin projector is needed: the two outer structure maps retain the required projection. Equation (112) removes the interior inverse/forward exterior maps and repeated expansion and contraction. It retains the same two Riesz maps algebraically; it does not replace covectors by coefficient vectors. Only one carrier workspace is live, with the original right-hand-side chunk bound. G is factored by spatial occupation blocks, not as a global dense structure matrix.

The normalization and selected-state constraints do not restore the need for repeated metric solves. If the columns of C are the selected states, put $K = C^{\mathsf{T}}MC$ and define

$$A = M^{-1} - CK^{-1}C^{\mathsf{T}}, \quad P_R = I - MCK^{-1}C^{\mathsf{T}}, \quad P_L = P_R^{\mathsf{T}}. \tag{113}$$

The exact constrained inverse preconditioner is

$$AF^*D^{-1}FA = P_L\left[M^{-1}F^*D^{-1}FM^{-1}\right]P_R. \tag{114}$$

Only the already available MC and the small state Gram K are needed for the two outer corrections. Thus two coordinate-metric CG solves in the composite evaluation are unnecessary. Multiplication by

the independent horizontal coordinate basis before and after this expression gives the same constrained operator. The identity applies columnwise when different roots use different shifted carrier diagonals.

The same closure identity permits direct evaluation of the complete-sector transport as

$$\mathbf{C} = G^{-1}\mathcal{W}^{*}\mathcal{T}(p)^{-1}b. \tag{115}$$

This does not license the corresponding replacement in Eq. (109) for a nonclosed selection. Moreover, CI transport used only as an initial guess for a later fully relaxed root solve is distinct from the transport used to define a keyframe objective: altering the former can change iteration cost and root tracking, while altering the latter also changes the finite model and its pullback. These two roles must not be conflated.

10.1 Complete-sector CI coordinates

For a complete spin sector, the orbital map need not appear in any inner CI iteration. Factor the orbital-independent spin Gram by occupation blocks, $G = LL^{\mathsf{T}}$, and put

$$U = \mathcal{W}L^{-\mathsf{T}}, \qquad U^{\mathsf{T}}U = I, \qquad UU^{\mathsf{T}} = P. \tag{116}$$

Here L is block diagonal in spatial occupations; no global dense factor is formed. The square change of selected-space coordinates and its inverse are

$$J = U^{\mathsf{T}}\mathcal{T}\mathcal{W}, \qquad J^{-1} = G^{-1}\mathcal{W}^{\mathsf{T}}\mathcal{T}^{-1}U. \tag{117}$$

Indeed, $UJ = \mathcal{T}\mathcal{W}$ follows from $P\mathcal{T}\mathcal{W} = \mathcal{T}\mathcal{W}$. Consequently, the original generalized eigenproblem is exactly equivalent to

$$H\mathbf{C} = M\mathbf{C}E \quad \Longleftrightarrow \quad (U^{\mathsf{T}}\mathcal{H}U)\mathbf{x} = \mathbf{x}E, \qquad \mathbf{x} = J\mathbf{C}. \tag{118}$$

The dimension remains the number of independent selected spin-adapted structures, not the dimension of the ambient spin-string product carrier. An inner Hamiltonian action is $\mathbf{x} \mapsto U^{\mathsf{T}}\mathcal{H}U\mathbf{x}$; its overlap action is the identity. Thus no inner exterior orbital map or orbital-dependent overlap action is mathematically needed. In the shared Davidson solver, the overlap images of the trial basis are then the trial basis itself. Their third stored basis matrix, associated matrix–vector products and separate reorthogonalization updates are omitted. The original subspace capacity is retained rather than spending the freed memory on a larger subspace.

The existing complete-sector inverse preconditioner transforms by congruence to $U^{\mathsf{T}}D^{-1}U$. In particular, $r_C = J^{\mathsf{T}}r_x$ and $JQ_CJ^{\mathsf{T}} = U^{\mathsf{T}}D^{-1}U$ imply that the preconditioned correction is the same physical direction. Metric orthogonalization of the original trial vectors is ordinary orthogonalization of their images under J. The change of coordinates therefore does not replace Davidson by a different physical subspace algorithm. For a pure target-spin residual, its Euclidean norm in these coordinates is its physical carrier norm. The same requested physical stationarity accuracy is retained; a new relaxed threshold is not introduced. The algebraic backward-error gates are evaluated in the isometric coordinates with the same numerical tolerances. Those gates are coordinate dependent, so the last accepted CI iteration need not be identical to that in the original coefficient representation; the physical stationarity norm is invariant. Initial trial vectors must first be mapped with J, rather than reinterpreted as coordinates in a different basis. At the exit, return

$$\mathbf{C} = J^{-1}\mathbf{x}, \qquad HC = J^{\mathsf{T}}h, \qquad MC = J^{\mathsf{T}}\mathbf{x}, \qquad h = (U^{\mathsf{T}}\mathcal{H}U)\mathbf{x}. \tag{119}$$

The retained Ritz images supply h without another Hamiltonian application. Only the boundary maps use $\mathcal{T}$ or its inverse. Carrier arrays are processed in budgeted RHS chunks and the sigma contraction retains its bounded internal tiles. No additional carrier-sized array is stored for each thread. This elimination requires the complete-sector identity; it is not valid for a general nonclosed selected VB space.

10.2 Cross-overlap generator moments

CI transport also requires overlaps between two orbital representations. Write $T_s = T_{s\alpha} \otimes T_{s\beta}$, where the exterior maps on side $s \in \{l, r\}$ are induced by the same spatial upper-triangular matrix R_s for both spins. For arbitrary selected coefficients, define

$$L = T_l \mathcal{W} c_l, \qquad R = T_r \mathcal{W} c_r, \qquad f = L^{\mathsf{T}} R. \tag{120}$$

These physical amplitudes lie in the ambient spin-string product carrier; no closure of range($\mathcal{W}$) under T_s is assumed. The cross-overlap action can be evaluated in the exact order

$$c \longmapsto \mathcal{W}^{\mathsf{T}} T_l^{\mathsf{T}} [T_r \mathcal{W} c]. \tag{121}$$

Each exterior map and its directional derivative is applied once to the shared amplitude array. In particular, it is unnecessary to reconstruct every output product from all source columns after the transforms.

The adjoint admits an additional contraction reduction. Let $\hat{E}_{pq} = \hat{a}^{\dagger}_{p\alpha} \hat{a}_{q\alpha} + \hat{a}^{\dagger}_{p\beta} \hat{a}_{q\beta}$ and define

$$D_{pq} = \langle R | \hat{E}_{pq} | L \rangle, \qquad K_s = \delta R_s R_s^{-1}. \tag{122}$$

The representation derivatives are $\delta L = \hat{E}(K_l) L$ and $\delta R = \hat{E}(K_r) R$, with $\hat{E}(K) = \sum_{pq} K_{pq} \hat{E}_{pq}$. Thus the upper-triangular orbital adjoints are

$$g_l = \mathrm{upper}(D R_l^{-T}), \qquad g_r = \mathrm{upper}(D^{\mathsf{T}} R_r^{-T}). \tag{123}$$

Here upper retains the diagonal and upper triangle; the transpose on D follows from real one-body matrix elements, not from equality of the left and right states. For a simultaneous direction,

$$\delta D_{pq} = \langle \delta R | \hat{E}_{pq} | L \rangle + \langle R | \hat{E}_{pq} | \delta L \rangle, \tag{124}$$

$$\delta g_l = \mathrm{upper}\big[\delta D R_l^{-T} - D R_l^{-T} (\delta R_l)^{\mathsf{T}} R_l^{-T}\big], \tag{125}$$

$$\delta g_r = \mathrm{upper}\big[(\delta D)^{\mathsf{T}} R_r^{-T} - D^{\mathsf{T}} R_r^{-T} (\delta R_r)^{\mathsf{T}} R_r^{-T}\big]. \tag{126}$$

Only one-body Slater–Condon links and small orbital triangular solves are needed. No Hamiltonian sigma, cofactor, or extra structure projection is required for these adjoints. Residual-weighted transport curvature uses these identities without setting the finite CI Ritz residual to zero.

If A and B denote the numbers of spin strings and r the number of source columns, the source-factor adjoint includes reconstruction work $O(|\Omega| r)$, in addition to its exterior transforms and generator contractions. For broad touched support, $|\Omega| = AB$ and $r = B$. Equations (121)–(126) replace that reconstruction by $O(\mathrm{nnz}\,\mathcal{W} + n_{\mathrm{act}}^2 AB)$ work. This is a different exact contraction order, not a low-rank truncation. It requires one or two shared carrier arrays for the overlap action, and two to four for its adjoint jet, depending on which side varies. They are admitted only within the existing caller workspace budget; there are no per-thread carrier copies. Narrow touched support or a carrier exceeding that budget retains the bounded source-panel order. Independent spin-dependent orbital maps also use that order, since Eq. (123) combines the two spins only for a common spatial chart. Neither schedule changes the selected VB space.

10.3 One final beta adjoint for image jets

Let Z be an orthogonal spin-string Hamiltonian or overlap image with A alpha rows and B beta columns. Its selected-space projection is

$$\mathcal{P}(Z) = \mathcal{W}^{\mathsf{T}} \,\mathrm{vec}(T_\alpha^{\mathsf{T}} Z T_\beta). \tag{127}$$

Here vec denotes the same fixed product ordering as the structure expansion. Its directional derivative is

$$\delta\mathcal{P} = \mathcal{W}^{\mathsf{T}} \,\mathrm{vec}(P T_\beta + Q\, \delta T_\beta), \tag{128}$$

$$P = T_\alpha^{\mathsf{T}} \delta Z + (\delta T_\alpha)^{\mathsf{T}} Z, \qquad Q = T_\alpha^{\mathsf{T}} Z. \tag{129}$$

These equations allow all alpha-pulled column contributions to accumulate before applying the beta adjoint. An ordinary image contributes to P; a left-representation derivative contributes to both P and Q. One ordered directional exterior sweep then returns $PT_\beta + Q\delta T_\beta$, after which the structure incidences are contracted once. The rule applies independently to each requested image, including the separate integral and representation forcing diagnostics. It does not assume that the selected structure space closes under T.

The source-panel alternative samples the transformed beta columns at the original touched products. Over all output panels of width t, this takes $O(|\Omega|B)$ beta reconstruction work, since B/t panels each evaluate length-t dots at $|\Omega|$ products. For broad support this is $O(AB^2)$. The final elementary exterior adjoint instead requires $O(n_{\mathrm{act}}^2 AB)$ work plus the structure incidence contraction. The older bounded group accumulator already avoids repeated incidence scans when resident; the reduction here concerns beta reconstruction, not an assertion that it always rescanned every structure per output tile.

This contraction order uses one shared $A \times B$ array per ordinary target and two per directional target. It replaces, rather than supplements, the group-projection payload and is admitted only when all target arrays fit its existing workspace quarter. Numerical residency and tile width are not increased. Narrow support or larger jets retain the exact bounded panel schedule. Integral sigma and adjoint sweeps, state weights and their spin-exchange multiplicities are unchanged; no approximate derivative or modified optimization tolerance is introduced.

10.4 Generated input jets without beta-row reconstruction

The input side admits the same ordering reduction. Reshape $\mathcal{W}c$ into $X \in \mathbb{R}^{A\times B}$ and an independent coefficient direction $\mathcal{W}q$ into δX. Then

$$Y = T_\alpha X T_\beta^{\mathsf{T}}, \tag{130}$$

$$\delta Y = (\delta T_\alpha) X T_\beta^{\mathsf{T}} + T_\alpha \delta X T_\beta^{\mathsf{T}} + T_\alpha X (\delta T_\beta)^{\mathsf{T}}. \tag{131}$$

Apply the ordered elementary exterior jet first on the alpha side, then on the beta side, to evaluate these expressions in two shared arrays. Generated sigma readers borrow their columns directly; they do not build another copy of each requested panel. The cost is $O(\mathrm{nnz}\,\mathcal{W} + n_{\mathrm{act}}^2 AB)$ instead of rebuilding every input beta panel in $O(|\Omega|B)$ work. For a coefficient-only direction with accepted Y already available, only δX is transformed; the accepted primal and residual panels remain shared with all consumers. For orbital directions the primal needed by the ordered jet is discarded after the derivative is formed when accepted panels already own Y.

The identity $\delta Y = \hat{E}(\delta R R^{-1})Y$ for a pure orbital direction is algebraically equivalent. The implementation uses Eq. (131), not an additional orbital inverse, so no new inverse-conditioned cancellation is introduced into the derivative. Its one or two arrays replace the original input-panel payload inside the same workspace quarter. Narrow support and arrays exceeding that budget use the existing exact generated-panel order. No assumption about closure of the selected structure space or any amplitude cutoff is needed.

10.5 Fixed selected metrics and singleton columns

Let $K_\sigma = T_{l,\sigma}^{\mathsf{T}} T_{r,\sigma}$, and label the touched original products by $g = (a_g, b_g) \in \Omega$. Suppose each touched β string occurs with only one α string. This single-entry-column property holds for a covalent occupation pattern: the β occupation is the complement of the α occupation. It does not require a complete structure space or spin-adapted closure. For b right-hand sides, first form the sparse structure expansion $D_{gj} =$

$\sum_I w_{gI} V_{Ij}$. The exact cross overlap is

$$P_{gh} = (K_\alpha)_{a_g a_h} (K_\beta)_{b_g b_h}, \tag{132}$$

$$Q_{gj} = \sum_{h \in \Omega} P_{gh} D_{hj}, \qquad Z_{Ij} = \sum_{g \in \Omega} w_{gI} Q_{gj}. \tag{133}$$

For a source tile $J \subset \Omega$, construct only the columns $A_J = K_\alpha[e_{a_h}]_{h \in J}$ and $B_J = K_\beta[e_{b_h}]_{h \in J}$. Their coefficient-free basis images are shared by all right-hand sides. Sample $(P_J)_{gh} = (A_J)_{a_g h} (B_J)_{b_g h}$ and accumulate $Q \leftarrow Q + P_J D_J$. The directional version uses precisely

$$\dot{P}_{gh} = (\dot{K}_\alpha)_{a_g a_h} (K_\beta)_{b_g b_h} + (K_\alpha)_{a_g a_h} (\dot{K}_\beta)_{b_g b_h}. \tag{134}$$

The ordered exterior jets already provide both factor derivatives. Unlike transforming coefficient-weighted source factors independently for each right-hand side, the basis-image work is no longer multiplied by b. Writing $t = |\Omega|$ and s_σ for the number of one-spin shear edges, the arithmetic is

$$O\Big(t(s_\alpha + s_\beta) + t^2 b + \mathrm{nnz}(\mathcal{W}) b\Big), \tag{135}$$

with tile width j and bounded RHS width b requiring $O((U_\alpha + U_\beta + t) j + tb)$ workspace. No $t \times t$ or $U_\alpha \times U_\beta$ numerical array is retained. This ordering is especially useful when projecting the selected metric, rather than repeatedly mapping each selected basis vector through the same exterior operators.

At a fixed chart, the selected metric

$$M = \mathcal{W}^* (K_\alpha \otimes K_\beta) \mathcal{W} \tag{136}$$

is immutable. If its N_{str}^2 entries fit the existing workspace, project M once with the tiled multi-RHS equations above and reuse Mv in subsequent metric CG iterations. This removes repeated exterior evaluations, not the selected-space projection or the metric. The same iterative solve, initial vector, preconditioner, and residual criterion are retained. A larger selected metric remains an exact tiled operator. Storage admission changes only the representation of this fixed linear operator; it never excludes a wave function from orthogonal CI. The implementation partitions the existing metric workspace into one half for construction and one quarter each for the selected identity and result; it introduces no string-pair storage and no larger memory allowance.

10.6 Metric-aware Davidson corrections in a selected space

The dimensions of the selected CI problem and the orthogonal contraction carrier are distinct. Davidson solves

$$H\mathbf{C} = EM\mathbf{C}, \qquad \mathbf{C}^{\mathsf{T}} M \mathbf{C} = 1, \tag{137}$$

in the original selected structure space, not in the string-product space. Let $F = \mathcal{T}\mathcal{W}$, so $H = F^* \mathcal{H} F$ and $M = F^* F$. A Ritz residual $\mathbf{r} = H\mathbf{C} - EM\mathbf{C}$ is a covector on that coefficient space. The metric representative of this covector is

$$\mathbf{z} = M^{-1} \mathbf{r}. \tag{138}$$

If the current Davidson basis is V, the Galerkin condition $V^{\mathsf{T}} \mathbf{r} = 0$ immediately gives

$$V^{\mathsf{T}} M \mathbf{z} = V^{\mathsf{T}} \mathbf{r} = 0. \tag{139}$$

Thus this correction lies outside the current subspace in the correct metric, up to the linear-solve and Ritz residual errors. Applying a raw coordinate diagonal to $\mathbf{r}$ does not have this property in general. In particular, diagonal normalization cannot remove correlations between nearly dependent Rumer columns.

For an invertible selected-basis change $\mathbf{C} = P\mathbf{x}$, the transformed residual and metric obey

$$\mathbf{r}_x = P^{\mathsf{T}} \mathbf{r}, \qquad M_x = P^{\mathsf{T}} M P, \qquad M_x^{-1} \mathbf{r}_x = P^{-1} M^{-1} \mathbf{r}. \tag{140}$$

The physical correction is unchanged. This covariance concerns the preconditioner; it does not assert that a coordinate-dependent finite stopping rule becomes invariant under every basis change.

Equation (138) is a metric preconditioner, not an exact inverse of $H - EM$. It removes the coefficient-space geometry but does not model every physical excitation gap. For a selected metric retained within the existing budget, factor $M = LL^{\mathsf{T}}$ once and use two triangular solves. The selected inverse is essential: for a nonclosed selection it cannot be replaced by the selected principal block of a complete-space inverse. No Hessian, structure Hamiltonian, full string-pair Gram or carrier-vector preconditioner is assembled. The original structure coefficients remain the Davidson variables, and all accuracy budgets are retained.

Once this fixed selected M is available, repeating its exterior back-projection during every Hamiltonian action is also unnecessary. For an input block V, evaluate only HV through generated product tiles and return MV from the small selected matrix. This removes the overlap image's alpha adjoint, beta sampling and structure projection from the product sweep. The Hamiltonian projection, its exact support and its Slater–Condon links are unchanged. The ordinary factor sweep retains its joint H/M contraction. Larger selected metrics continue to use the same matrix-free generalized problem with the existing bounded preconditioner; the storage decision never changes the variational space or disables orthogonal CI.

11 Correspondence with the current implementation

The original column contraction corresponds to revision `a5004add`; Section 3.2 adds the minimum mixed-support decomposition. Paths below are relative to `src/`; this table records the implementation boundary rather than introducing physical assumptions.

Mathematical operation	Implementation
Original support and minimum source cover	`vbscf/structures/assembly/sources.cpp`
VB source generation, sampling and $\mathcal{W}^*$	`vbscf/structures/assembly/source_tiles.cpp` `apply_source_tiles()`
Common factor sweep and integral adjoint	`mcscf/ci/factors.cpp`
Selected-structure response and its adjoints	`vbscf/structures/assembly/source_response.cpp`
Source-state orbital directions and HVP pullback	`vbscf/structures/assembly/source_direction.cpp`
Broad-support accepted/tangent sweep	`vbscf/structures/assembly/product_jet.cpp`
Shared H-image projection and generator adjoint	`vbscf/structures/assembly/source_image.cpp`
Generator and residual-weighted factor adjoints	`mcscf/ci/factor_generator.cpp`
Joint source/channel work plan	`mcscf/ci/factor_plan.cpp`
Hamiltonian channel kernel, Eq. (23)	`mcscf/ci/sigma.cpp` `build_hamiltonian_pair_kernel()`
Slater–Condon spin connections	`mcscf/ci/spin_channels.cpp`
Exterior transforms and causal adjoint rows	`vbscf/structures/orthogonal_ci/exterior_transform.cpp` `projection()` and `apply_adjoint_left_block()`
Independent complete-carrier comparisons	`../tests/vbscf/unit/test_structure_tiles.cpp`

The active-space forward entries complete the computational fixed-spin string tables before constructing the action. Original input incidences are remapped into those tables; neither the selected VB structures nor their independent coefficients are enlarged. This completion supplies the exterior maps and Slater–Condon links without allocating numerical pair tables or product amplitudes. It must occur in the production entries, not only in selected-space test fixtures.

Fixed-point structure Hamiltonian/overlap actions choose between the common factor sweep and generated product tiles. Accepted-state integral gradients choose between the exact factor and product contractions by their arithmetic costs. Raw state coefficients are retained in structure space. Broad product-order accepted, directional and reverse consumers use Section 9. Compact accepted wave-function/residual panels may remain under one shared residency cap; unbounded accepted and directional carrier fields are deleted. A selected-space eigenvector need not be stationary on the full carrier: its residual contribution to the exterior-transform gradient is retained. Accepted-state transition densities use Eq. (68) when advantageous. The residual-panel kernel of Eq. (70) is available and independently tested; source-order accepted states now omit accepted orthogonal and residual arrays. The coupled coefficient-response consumer in Section 8.5 uses neither accepted nor response product arrays when its factor contraction order is cheaper. Section 8.7 extends that ownership to source-state orbital directions and HVPs. Broad product-order derivative consumers now share the generated product jet. Complete-space inverse/preconditioner consumers remain a separate ownership boundary.

Main result. Equations (36)–(42) give the exact selected-space restriction through spin factors and Slater–Condon actions, without a closure assumption or a complete transformed product array.